

Bitonic Sorting

1 Parallel Merge Sort

1.1 Serial Merge Sort

The runtime for serial merge sort using divide and conquer is $O(n \log n)$, which has been proved [here](#).

1.2 Parallel Merge Sort

First explore some special cases:

1. Divide into 2 parts, when merge need to traverse n :

$$T(n, p) = T\left(\frac{n}{2}, \frac{p}{2}\right) + O(n) \quad (1)$$

2. Divide into 4 parts, the first merge needs to traverse $\frac{n}{2}$, the second merge needs to traverse n :

$$T(n, p) = T\left(\frac{n}{4}, \frac{p}{4}\right) + O\left(\frac{n}{2}\right) + O(n) \quad (2)$$

Now for p processors, divide into p parts:

$$T(n, p) = T\left(\frac{n}{p}, 1\right) + O\left(n + \frac{n}{2} + \frac{n}{4} + \dots + \frac{n}{p/2}\right) \quad (3)$$

Recall the serial merge sort runtime:

$$T(n) = T(n, 1) = O(n \log n) \quad (4)$$

Therefore we have:

$$T(n, p) = O\left(\frac{n}{p} \log \frac{n}{p} + n\right) \quad (5)$$

To use resources optimally, the number of processors should match the level of parallelism that the algorithm can effectively exploit. For parallel merge sort, this means having a number of processors that aligns with the depth of the recursion tree, which is $\log n$. Using more processors than this would not make the algorithm run faster because the bottleneck becomes the merging process, which has to wait for all sub-problems at a given level to be completed before it can proceed to the next level of merging.

2 Bitonic Sort

2.1 Overview

Let l_1 and l_2 be non-decreasing sorted lists, including:

$$l_1 : x_0, x_1, \dots, x_{n-1} \quad (6)$$

$$l_2 : y_0, y_1, \dots, y_{n-1} \quad (7)$$

Then we can generate a minimum array and a maximum array:

$$l_{min} : \min(x_0, y_{n-1}), \min(x_1, y_{n-2}), \dots, \min(x_{n-1}, y_0) \quad (8)$$

$$l_{max} : \max(x_0, y_{n-1}), \max(x_1, y_{n-2}), \dots, \max(x_{n-1}, y_0) \quad (9)$$

An example is shown below:

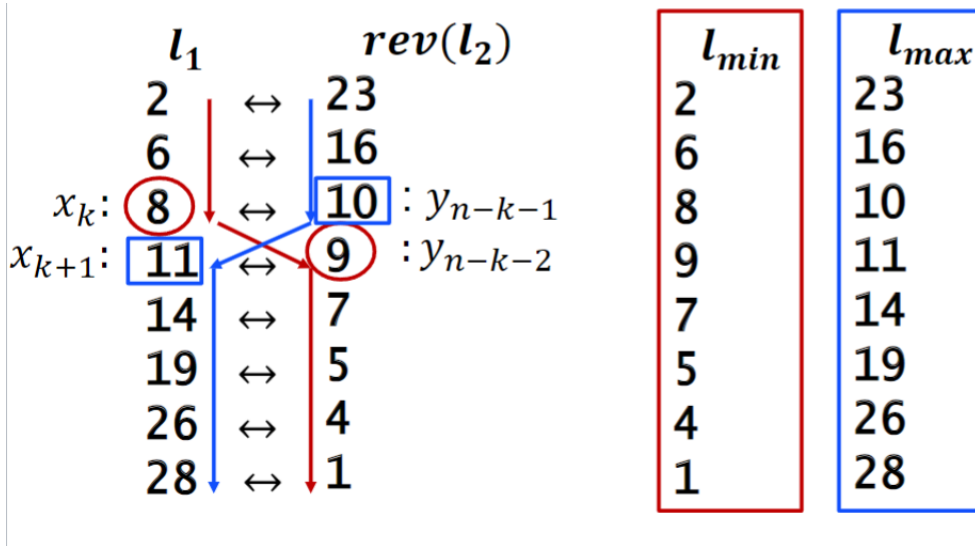


Figure 1: Bitonic Sort Example

Based on the observation, we can see that:

$$\max(l_{min}) \leq \min(l_{max}) \quad (10)$$

This could be proven. Assume the transition index as k in l_1 , then:

$$\max(l_{min}) = \max(x_k, y_{n-k-2}) \quad (11)$$

And:

$$\min(l_{max}) = \min(x_{k+1}, y_{n-k-1}) \quad (12)$$

The proof steps include:

1. Because l_1 is sorted, so $x_k \leq x_{k+1}$.

2. Because the transition is after k , so $x_k \leq y_{n-k-1}$
3. Also because the transition is after k , so $y_{n-k-2} \leq x_{k+1}$
4. Because l_2 is sorted, so $y_{n-k-2} \leq y_{n-k-1}$

Therefore, we can easily see that $\max(l_{min}) \leq \min(l_{max})$.

2.2 Runtime

The runtime for bitonic sort is $O(n \log \frac{n}{p})$, and the proof is shown below:

$$\begin{aligned}
 & \bullet T(n, p) = T\left(\frac{n}{2}, \frac{p}{2}\right) + O\left(\frac{n}{p}\right) + T\left(\frac{n}{2}, \frac{p}{2}\right) \\
 & \bullet T(n, p) = 2T\left(\frac{n}{2}, \frac{p}{2}\right) + O\left(\frac{n}{p}\right) \\
 & \bullet T(n, p) = 2\left(2T\left(\frac{n}{4}, \frac{p}{4}\right) + O\left(\frac{n/2}{p/2}\right)\right) + O\left(\frac{n}{p}\right) \\
 & \bullet T(n, p) = 2^2 T\left(\frac{n}{4}, \frac{p}{4}\right) + O\left(2\frac{n}{p} + \frac{n}{p}\right) \\
 & \bullet T(n, p) = 2^3 T\left(\frac{n}{8}, \frac{p}{8}\right) + O\left(2^2 \frac{n}{p} + 2\frac{n}{p} + \frac{n}{p}\right) \\
 & \bullet T(n, p) = pT\left(\frac{n}{p}, 1\right) + O\left(\frac{n}{p} \left(1 + 2 + 2^2 + \dots + \frac{p}{2}\right)\right) \quad (p-1) \\
 & \bullet \text{Each processor sorts locally in } O\left(\frac{n}{p} \log \frac{n}{p}\right) \text{ time.} \\
 & \bullet T(n, p) = O\left(\cancel{\frac{n}{p}} \log \frac{n}{p} + n\right) = O\left(n \log \frac{n}{p}\right)
 \end{aligned}$$

Figure 2: Bitonic Sort Runtime Proof

2.3 Bitonic Sequence

We call x_0, x_1, \dots, x_{n-1} as bitonic sequence if there is an index k such that (could be one of the three):

1. x_0, x_1, \dots, x_k is non-decreasing and $x_{k+1}, x_{k+2}, \dots, x_{n-1}$ is non-increasing.
2. x_0, x_1, \dots, x_k is non-increasing and $x_{k+1}, x_{k+2}, \dots, x_{n-1}$ is non-decreasing.
3. There is a cyclical shift of the sequence that makes the first two situations true.

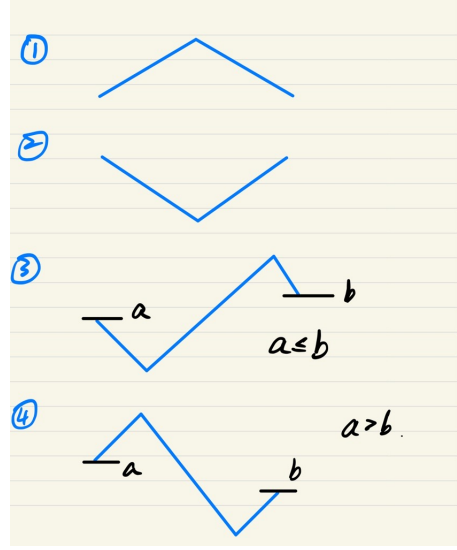


Figure 3: Bitonic Sequence

2.4 Bitonic Split

2.4.1 Definition

Bitonic split of a **bitonic sequence** $l = x_0, x_1, \dots, x_{n-1}$ is defined as decomposition of l into:

$$l_{\min} = \min(x_0, x_{\frac{n}{2}}), \min(x_1, x_{\frac{n}{2}+1}), \dots, \min(x_{\frac{n}{2}-1}, x_{n-1}) \quad (13)$$

$$l_{\max} = \max(x_0, x_{\frac{n}{2}}), \max(x_1, x_{\frac{n}{2}+1}), \dots, \max(x_{\frac{n}{2}-1}, x_{n-1}) \quad (14)$$

2.4.2 Lemma

Let l be a bitonic sequence, and l_{\min} and l_{\max} result from the bitonic split, then we have:

1. l_{\min} and l_{\max} are bitonic.
2. $\max(l_{\min}) \leq \min(l_{\max})$

The proof is shown below. if l is the type T1 bitonic sequence, then:

1. When $x_0 > x_{\frac{n}{2}}$:

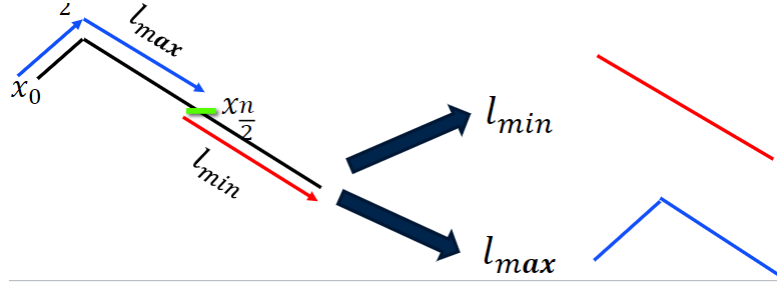


Figure 4: Case 2

2. When $x_0 > x_{\frac{n}{2}}$:

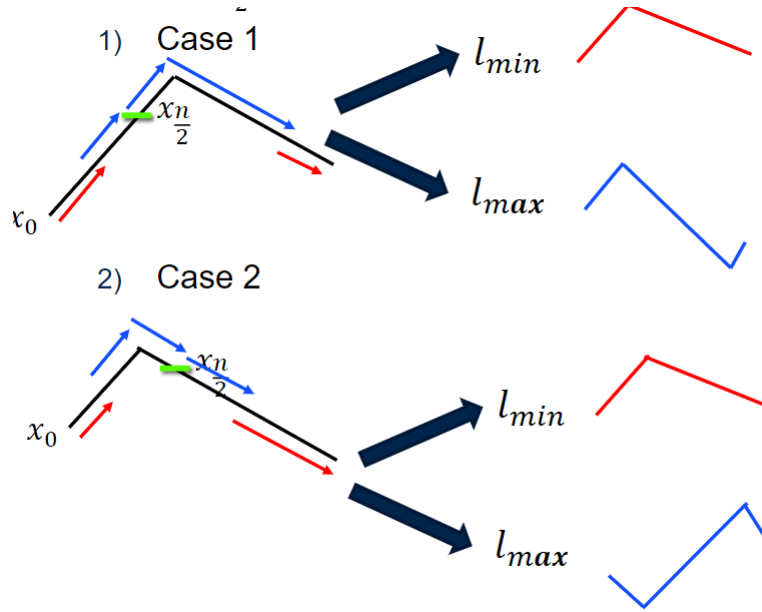


Figure 5: Case 2

2.4.3 Sub-Lemma

Let l be a bitonic sequence and l' be a cyclical shift. Assume bitonic split l into l_{min} and l_{max} and bitonic split l' into l'_{min} and l'_{max} , then we have:

1. l'_{min} is a cyclical shift of l_{min}
2. l'_{max} is a cyclical shift of l_{max}

So the cyclical shift does not change the bitonic nature of the sequence nor the min(or max) element in the sequence.

2.5 Bitonic Merge

Bitonic merge means turning a bitonic sequence into a sorted sequence using repeated bitonic split operations. The runtime could be expressed as:

$$BM(p, p) = O(\log p) \quad (15)$$

2.6 Bitonic Sort Example

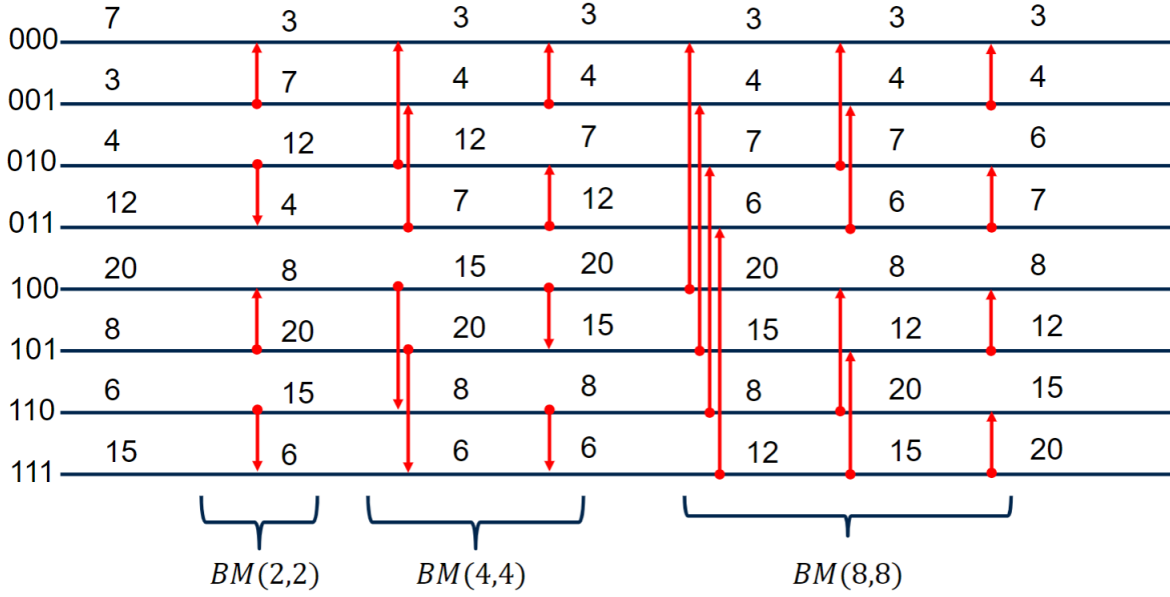


Figure 6: Bitonic Sort Example

Assume we have processors p_1, p_2, \dots, p_8 and the corresponding value 7, 3, 4, 12, 20, 8, 6, 15. Then we first divide these values into small sets with only 2 values. Then, we perform **bitonic merge** in each of the set. This includes first **bitonic split** the set and then combine them in sorted way. This is done by putting the split result into a specific processor order. Notice that here **in the first set of 2, we choose ascending order, which means $p_1 < p_2$, and in the second set of 2, we choose descending order, which means $p_3 > p_4$** . This is for the next level's merging. Because now we have only two values, bitonic split the set and put the results into specific order will be the sorting:

$$[p_1(7), p_2(3)] \rightarrow [l_{min}, l_{max}] \rightarrow [p_1(3), p_1(7)] \quad (16)$$

$$[p_3(4), p_4(12)] \rightarrow [l_{max}, l_{min}] \rightarrow [p_3(12), p_4(4)] \quad (17)$$

Because in this level we have two values in each set, so $BM(2, 2) = \log_2 2 = 1$. Now we move to the next level, with each set containing 4 values. Similar with before, **in the first set of 4, we choose ascending order, which means $p_1 < p_2 < p_3 < p_4$, and in the second set of 4, we choose descending order, which means $p_5 > p_6 > p_7 > p_8$** . Now we have:

$$[p_1(3), p_2(7), p_3(12), p_4(4)] \rightarrow [l_{min}, l_{max}] \rightarrow [p_1(3), p_2(4), p_3(12), p_4(7)] \quad (18)$$

$$[p_5(8), p_6(20), p_7(15), p_8(6)] \rightarrow [l_{max}, l_{min}] \rightarrow [p_5(15), p_6(20), p_3(8), p_4(4)] \quad (19)$$

But this is not enough for sorting the first set of 4. We still need to divide it into set of 2, then perform bitonic merge for it, for example:

$$[p_3(12), p_4(7)] \rightarrow [l_{min}, l_{max}] \rightarrow [p_3(7), p_4(12)] \quad (20)$$

$$[p_5(15), p_6(20)] \rightarrow [l_{max}, l_{min}] \rightarrow [p_5(20), p_4(15)] \quad (21)$$

This will take $\log_2 4 = 2$ steps. Similar with $BM(8, 8)$, which will take 3 steps. Therefore, the bitonic sort runtime (computation) could be expressed as:

$$BS(p, p) = BM(2, 2) + BM(4, 4) + \dots BM(p, p) \quad (22)$$

$$BS(p, p) = 1 + 2 + 3 + \dots + \log p = O(\log^2 p) \quad (23)$$

Because each step in this case need communications between processors, so the communication time is $O((\tau + \mu) \log^2 p)$.

If we have $n > p$, then we need to **sort n/p local elements before doing the bitonic sort, which will take $O(n \log n)$** , so:

$$\text{Computation T} = O\left(\frac{n}{p} \log \frac{n}{p} + \frac{n}{p} \log^2 p\right) \quad (24)$$

For the communication time, recall that μ is the time taken to send each unit of data. Here, each processor contains n/p data, and τ is the time for interaction between each processor, so:

$$\text{Communication T} = O\left(\tau \log^2 p + \mu \frac{n}{p} \log^2 p\right) \quad (25)$$