

# Parallel and Distributed Simulation

## 1 Introduction

### 1.1 Overview

Some definitions:

1. **Parallel Simulation:** It involves the execution of a **single** simulation on a collection of **tightly coupled processors** (such as a shared memory multiprocessor).

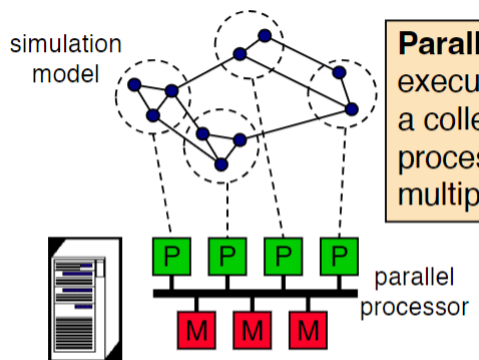


Figure 1: Parallel Simulation

2. **Replicated Trials:** It involves the execution of **several, independent but the same** simulation runs concurrently on different processors.




Figure 2: Replicated Trials

3. **Distributed Simulation:** It involves the execution of a **single** on a collection of **loosely** coupled processors, such as servers interconnected by a Local Area Network (**LAN**) and Wide Area Network (**WAN**).




Figure 3: Distributed Simulation

## 1.2 Distributed Simulation

There are several properties of distributed platforms:

1. **Private Memory:** Each processing node has its own private memory, so one node could not directly modify the memory of another
2. **Message-Passing:** Nodes communication by exchanging messages
3. **High Cost:** The cost of sending messages is high compared to local computation.

## 1.3 Reasons for Parallel/Distributed Simulation

- Requires too much time to complete a single simulation
- Requires rapid response for real-time decision making to manage operational systems
- Requires more memory than available on a single computer
- Requires geographically distributed people or resources

## 1.4 System and Time

Usually systems could be divided into:

1. **Physical system:** the actual or imagined system being modeled
2. **Simulation system:** a system that emulates the behavior of a physical system

And the times could be categorized into:

1. **Physical time:** time in the physical system
2. **Simulation (logic) time:** representation of physical time in the simulation
3. **Wallclock time:** time during the execution of the simulation, usually output from a hardware clock

## 2 Parallel Discrete Event Simulation

### 2.1 Model Setup

Recall the [airport example](#), first we transfer the physical process to **logical process**:

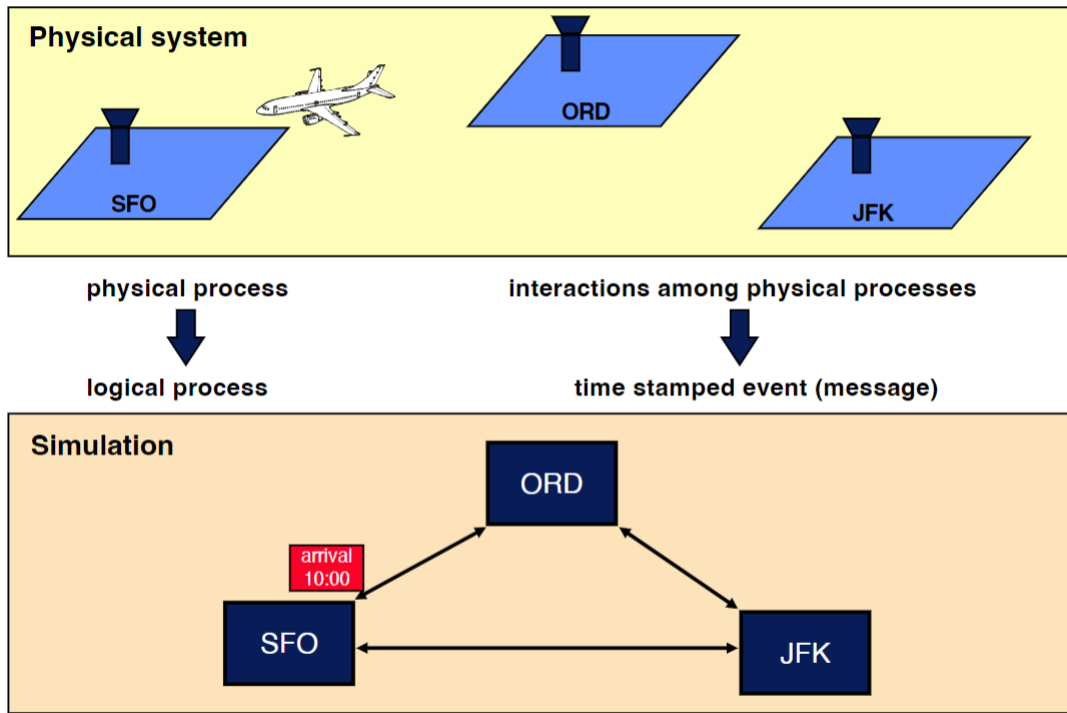


Figure 4: Physical to Logical Process

Some remarks:

1. Each airport simulator is a logical process (LP).
2. An LP can schedule events for other LPs by **sending messages**.
3. **No shared state** between LPs.
4. The **physical system** here is defined as the collection of interacting physical processes (airports).
5. The **simulation system** here is defined as:
  - A collection of logical processes (LPs)
  - Each LP models a physical process
  - The interactions between physical processes modeled by **scheduling events between LPs**

## 2.2 Parallel Implementation

To implement the simulation, we need to map LPs to different processors. It is allowed to have multiple LPs per processor. All the interactions will via messages. Some important rules:

## 2.3 Local Causality Constraint (LCC)

### 2.3.1 Definition

A model should be designed such that an event or state change at a given point in space and time can only affect, and be affected by, events and states in its **local neighborhood within a finite amount of time**. In other words, process incoming messages in time stamp order, including externally generated events.

### 2.3.2 Case Study

Assume we have several logical processes shown below:

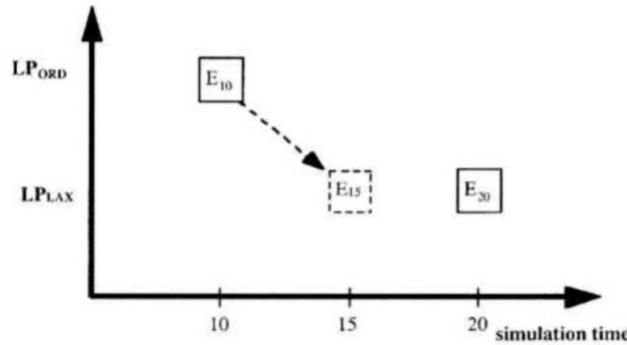


Figure 5: Airports Logical Process

Some clarifications:

1. ORD and LAX represent two local processing units (LPs) handling events.
2.  $E_{10}$  is an event in the queue at ORD with a timestamp of 10.  $E_{20}$  is an event in the queue at LAX with a timestamp of 20.
3. When ORD process event  $E_{10}$ , it may schedule a remote event  $E_{15}$  at LAX. **This means that  $E_{15}$  is queued at LAX but has a timestamp of 15, which is earlier than  $E_{20}$ .**

This is a good example of LCC. According to LCC, events must be processed in the order of their timestamps to ensure the simulation's accuracy. Therefore, **LAX could not process  $E_{20}$  before  $E_{15}$  because  $E_{15}$  might change the state or conditions that  $E_{20}$  depends on.** Therefore, LAX must wait to process  $E_{20}$  until after it processes  $E_{15}$ . This shows the challenge of **achieving concurrency** (the ability of multiple logical processes (LPs) to execute simultaneously) in distributed simulations.

### 3 Synchronization Problem

Based on the previous discussion, an algorithm is needed to ensure each LP processes events in time stamp order. Based on the observations:

1. **Ignoring simultaneous events:** This means we ignore the events that have identical timestamps. These events are often considered separately because they can be processed in any order without affecting the overall outcome.
2. **Adherence to the LCC:** This is sufficient to ensure that **the parallel simulation will produce the same results as a sequential execution**. If LCC is followed, the results will be consistent with a scenario where all events across all LPs are processed sequentially in timestamp order.

### 4 Conservative Synchronization

#### 4.1 Overview

Conservative synchronization avoids the risk of causality errors (events are processed out of their intended order, violating the natural cause-and-effect relationships that should be preserved) by **preventing any process from executing an event until it is certain that no other events with earlier timestamps can affect it**.

Some **key principles**:

1. **Blocking until safe (Change!!!):** Each processing unit (LP) blocks the execution of an event until it is guaranteed that no other event with an earlier timestamp can arrive.
2. **Lookahead:** This the minimum time into the future that an LP can predict events will not interfere. LPs can use lookahead to safely process events up to a certain point, knowing that earlier events will not occur.
  - **Link lookahead:** If an LP is at simulation time  $T$ , and **an outgoing link** has lookahead  $L_i$ , then any message sent on that link must have a timestamp of at least  $T + L_i$ .
  - **LP lookahead:** If an LP is at simulation time  $T$ , and has a lookahead  $L$ , then that LP must have a timestamp of at least  $T + L$ .
3. **Deadlock avoidance:** A common challenge in conservative synchronization is the potential for deadlock, **where all LPs are waiting for each other indefinitely**. Mechanisms such as **null messages** (messages that convey the absence of events within a certain time frame) are used to prevent deadlocks by providing LPs with the information needed to make progress.

The **advantages** of conservative synchronization include:

1. **Accuracy:** Guarantees that events are processed in the correct order, maintaining the accuracy and consistency of the simulation.

2. **Simplicity:** Conceptually straightforward and easier to implement compared to optimistic synchronization.

But there are also some **disadvantages**:

1. **Potential Inefficiency:** Can lead to significant idle time for LPs, waiting for the assurance that it is safe to proceed, which can reduce overall simulation performance.
2. **Scalability:** May not scale well for very large and complex simulations due to increased communication overhead and potential for blocking.

## 4.2 Chandy/Misra/Bryant (CMB) Null Message Algorithm

### 4.2.1 Introduction

In CMB algorithm, LPs periodically send special messages, called **null messages**, to their neighboring LPs. Null messages indicate that **no events with timestamps earlier than a certain value will be sent in the future**.

### 4.2.2 Assumptions

This algorithm needs several **assumptions**:

- LPs exchanging timestamped events (messages).
- Static network topology (LP-to-LP), no dynamic creation of LPs.
- Messages sent on each link are sent in timestamp order.
- Network provides reliable delivery, preserves order.

### 4.2.3 Procedures

The procedures include:

1. **Initial State:** Each LP starts with its local event queue containing events it needs to process.
2. **Sending Null Messages:** When an LP processes an event, it sends a null message to its neighboring LPs. The null message contains a timestamp **indicating the earliest time at which the LP might send a new event in the future**. This timestamp is typically **the current simulation time plus some lookahead value**, representing the minimum time into the future before which no new events will be generated.
3. **Receiving Null Messages:** When an LP receives a null message from a neighboring LP, it updates its understanding of the earliest possible future event from that neighbor. **This helps the receiving LP determine the safe time horizon up to which it can process its local events without waiting for earlier events from its neighbors.**

4. **Processing Events:** Each LP processes its local events in timestamp order up to the minimum of the null message timestamps received from all its neighbors.
5. **Deadlock Avoidance:** The algorithm uses null messages to avoid deadlock situations where all LPs are waiting indefinitely for events that might never arrive.

### 4.3 Time Creep Problem

Time creep occurs when LPs advance their local simulation clocks very slowly because they are overly cautious. Each LP waits for confirmation that it can safely process an event, often resulting in significant idle times as LPs wait for messages from other LPs. This conservatism can lead to the overall simulation progressing very slowly, as LPs make minimal progress due to the need for constant synchronization and assurance of event order.